Lesson 1.1: Why Distributed Computing?

DISTRIBUTED COMPUTING WITH SPARK SQL

Why Distributed Computing?

Brooke Wenig/Conor Murphy
Machine Learning Practice Lead/
Data Scientist, Databricks

UCDAVIS
Continuing and Professional Education

Copyright ©The Regents of the University of California

Slide 2: Learning Objectives

Learning Objectives

Motivate the business need for processing big data

Identify key concepts related to distributed computing

Copyright ©The Regents of the University of California

Slide 3: Qualities of Big Data

Qualities of Big Data

Volume

Velocity

Variety

Veracity

Copyright ©The Regents of the University of California

Slide 4: Big Data Defined

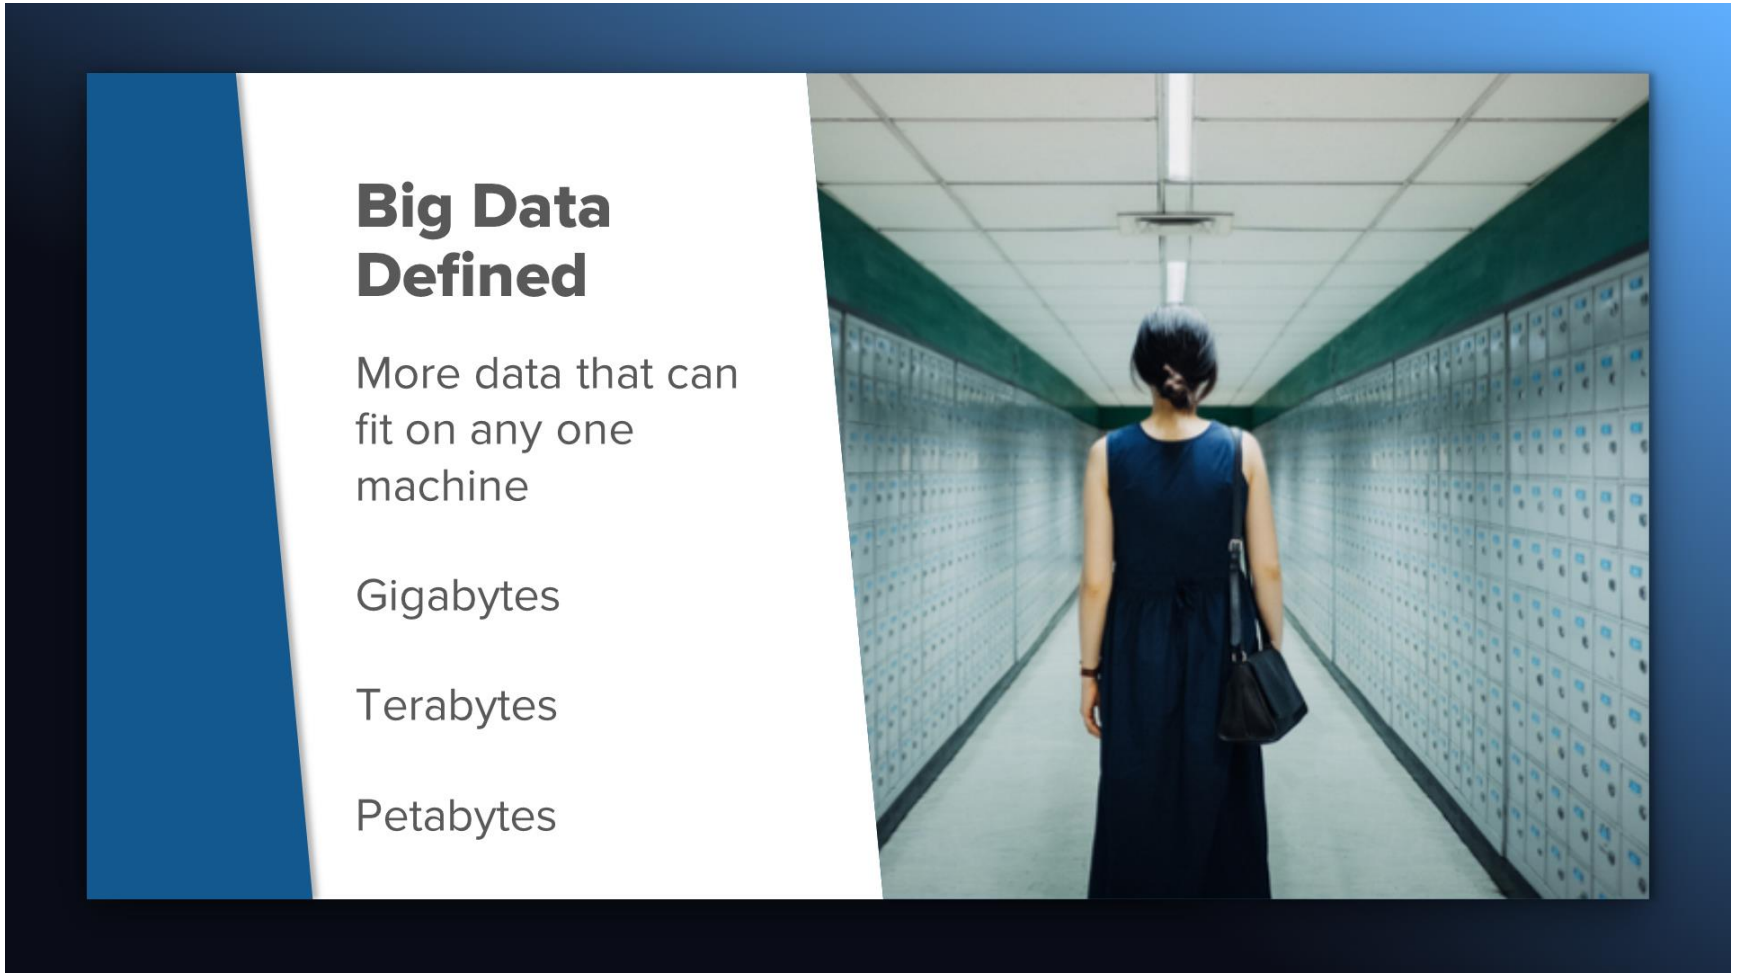

Copyright ©The Regents of the University of California

Slide 5: Apache Spark

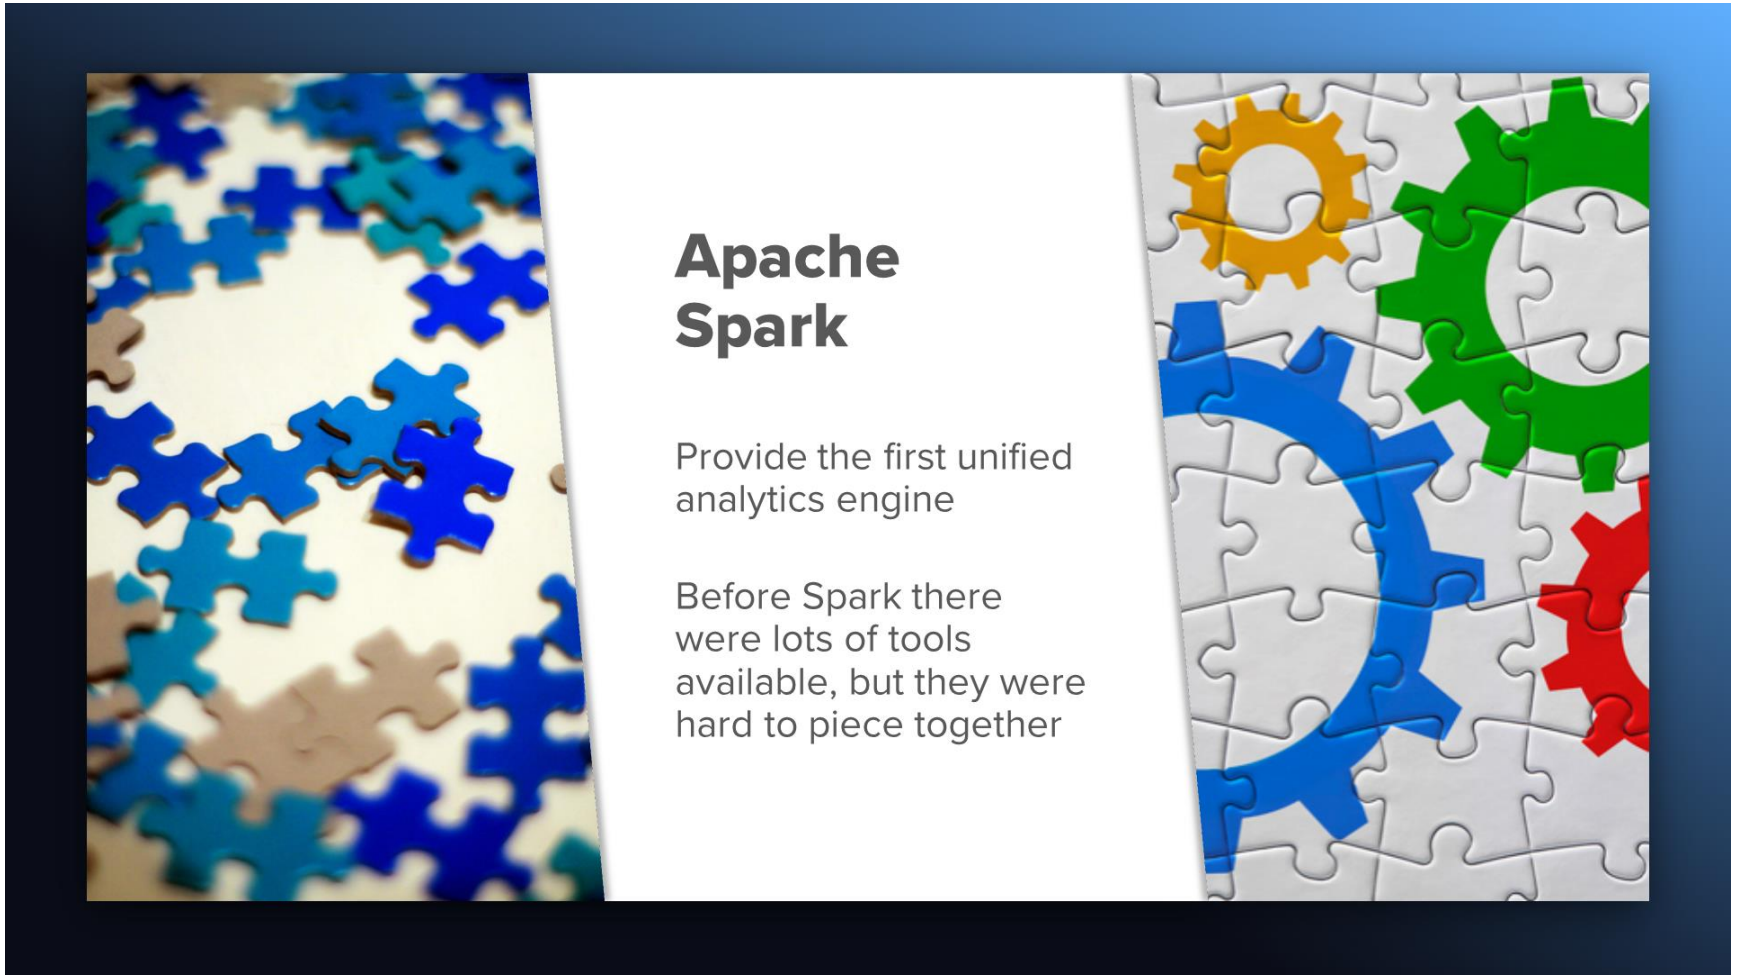

Copyright ©The Regents of the University of California

Slide 6: Spark is Multilingual and Supports Many Languages

Spark Supports Many Languages

SQL

Python

Scala

Java

R

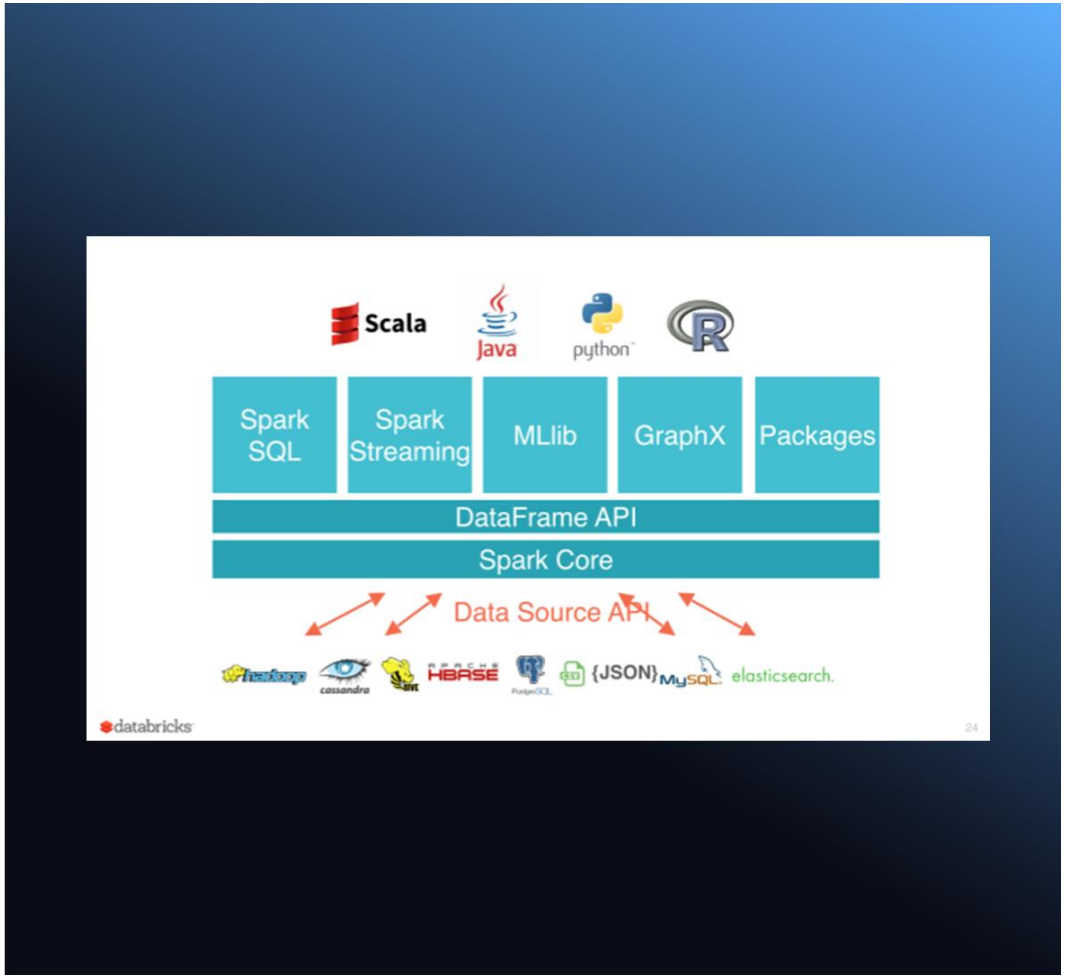

Copyright ©The Regents of the University of California

Slide 7: Why Spark is Popular

Why Spark is Popular

Reads & processes data from many sources

Works with many file types

Solves many data problems faced by analysts

Copyright ©The Regents of the University of California

Slide 8: Apache Spark: Origin Story

Copyright ©The Regents of the University of California

Slide 9: Apache Spark: Origin Story

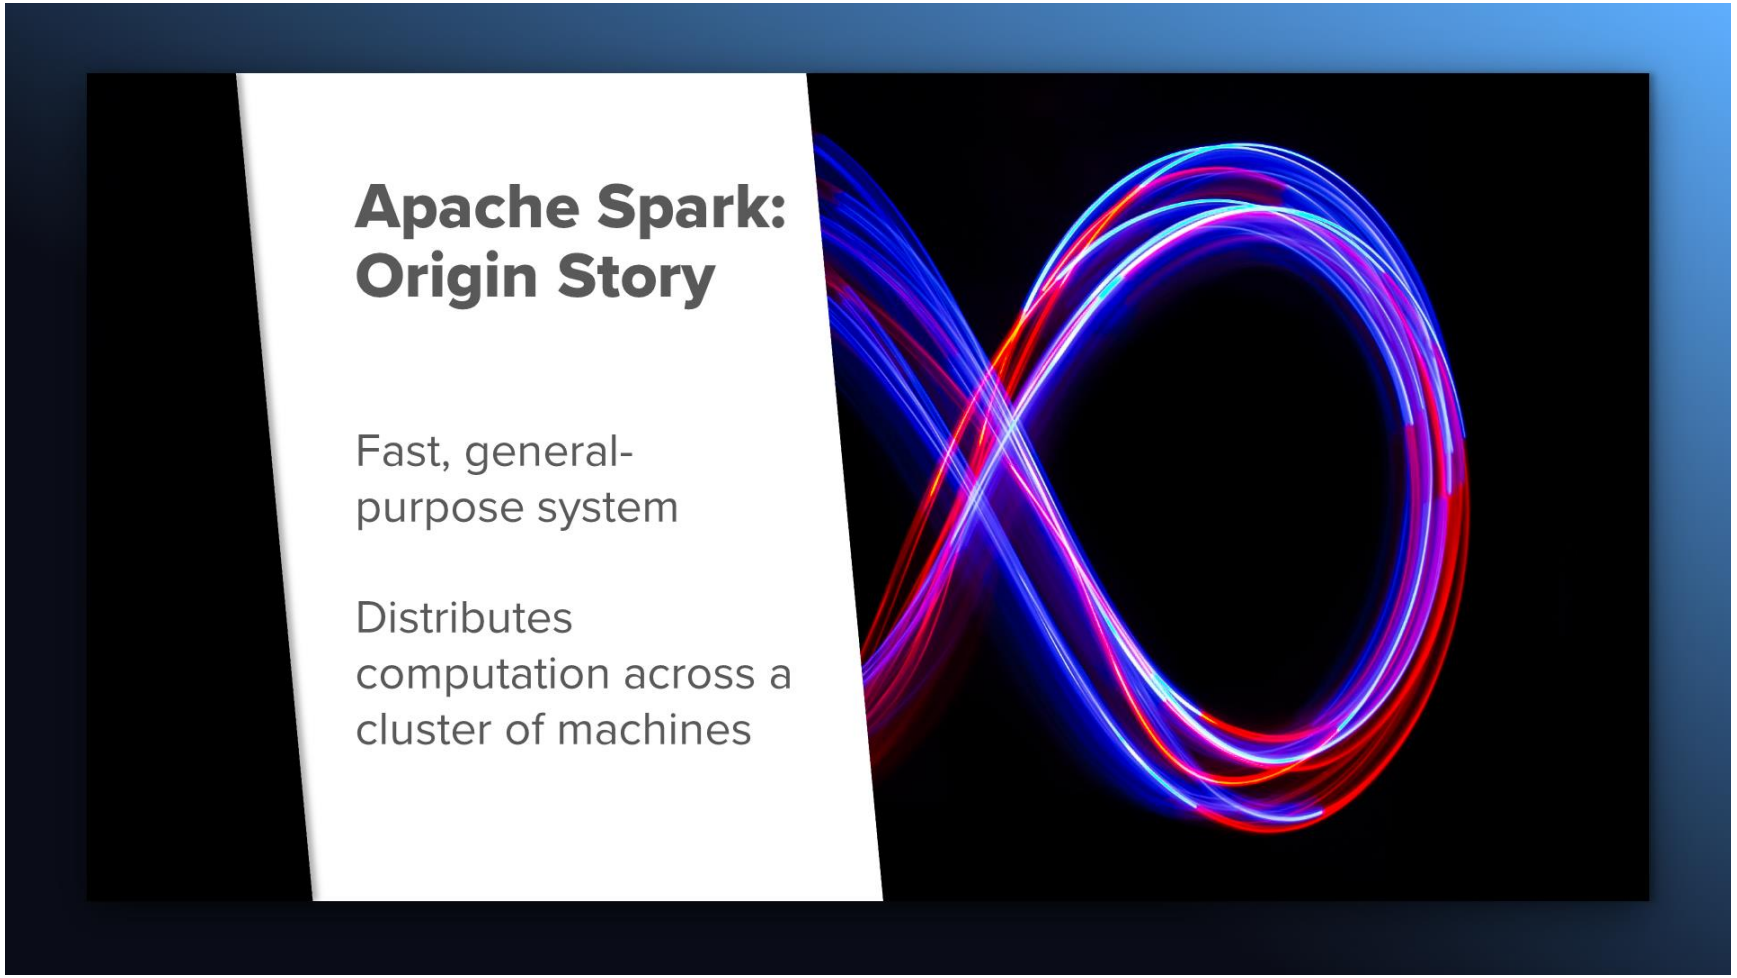

Copyright ©The Regents of the University of California

Slide 10: Apache Spark: Origin Story

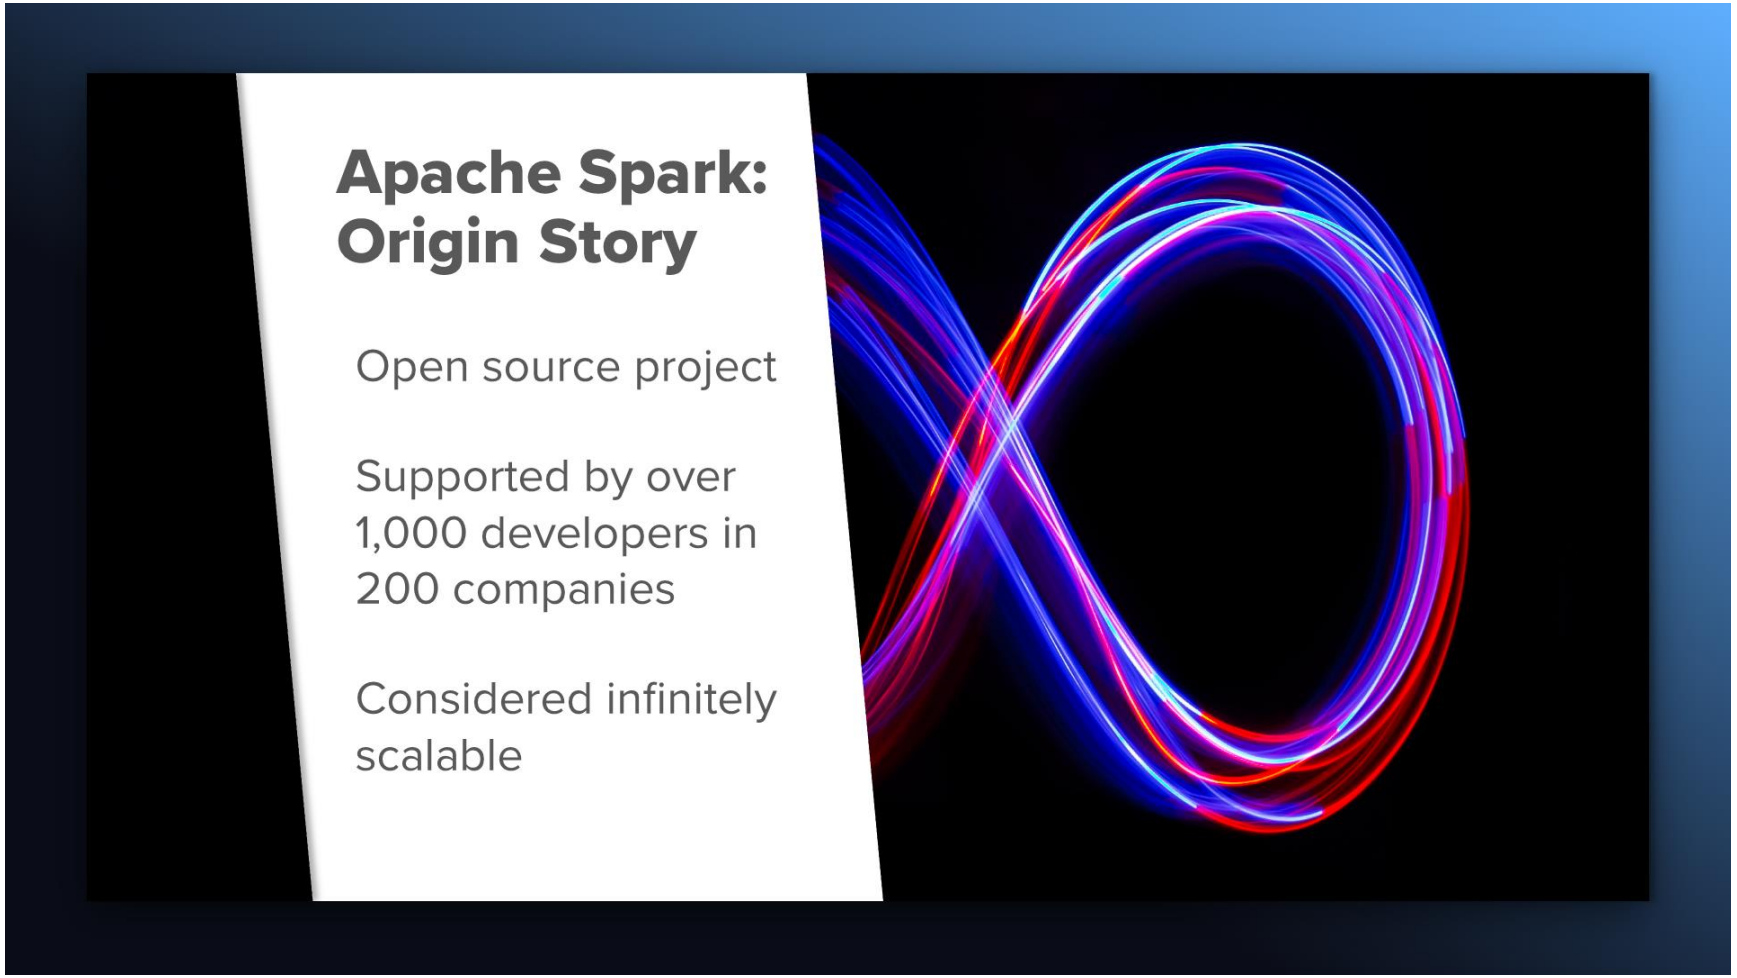

Copyright ©The Regents of the University of California

Slide 11: Let’s Count Some M&Ms

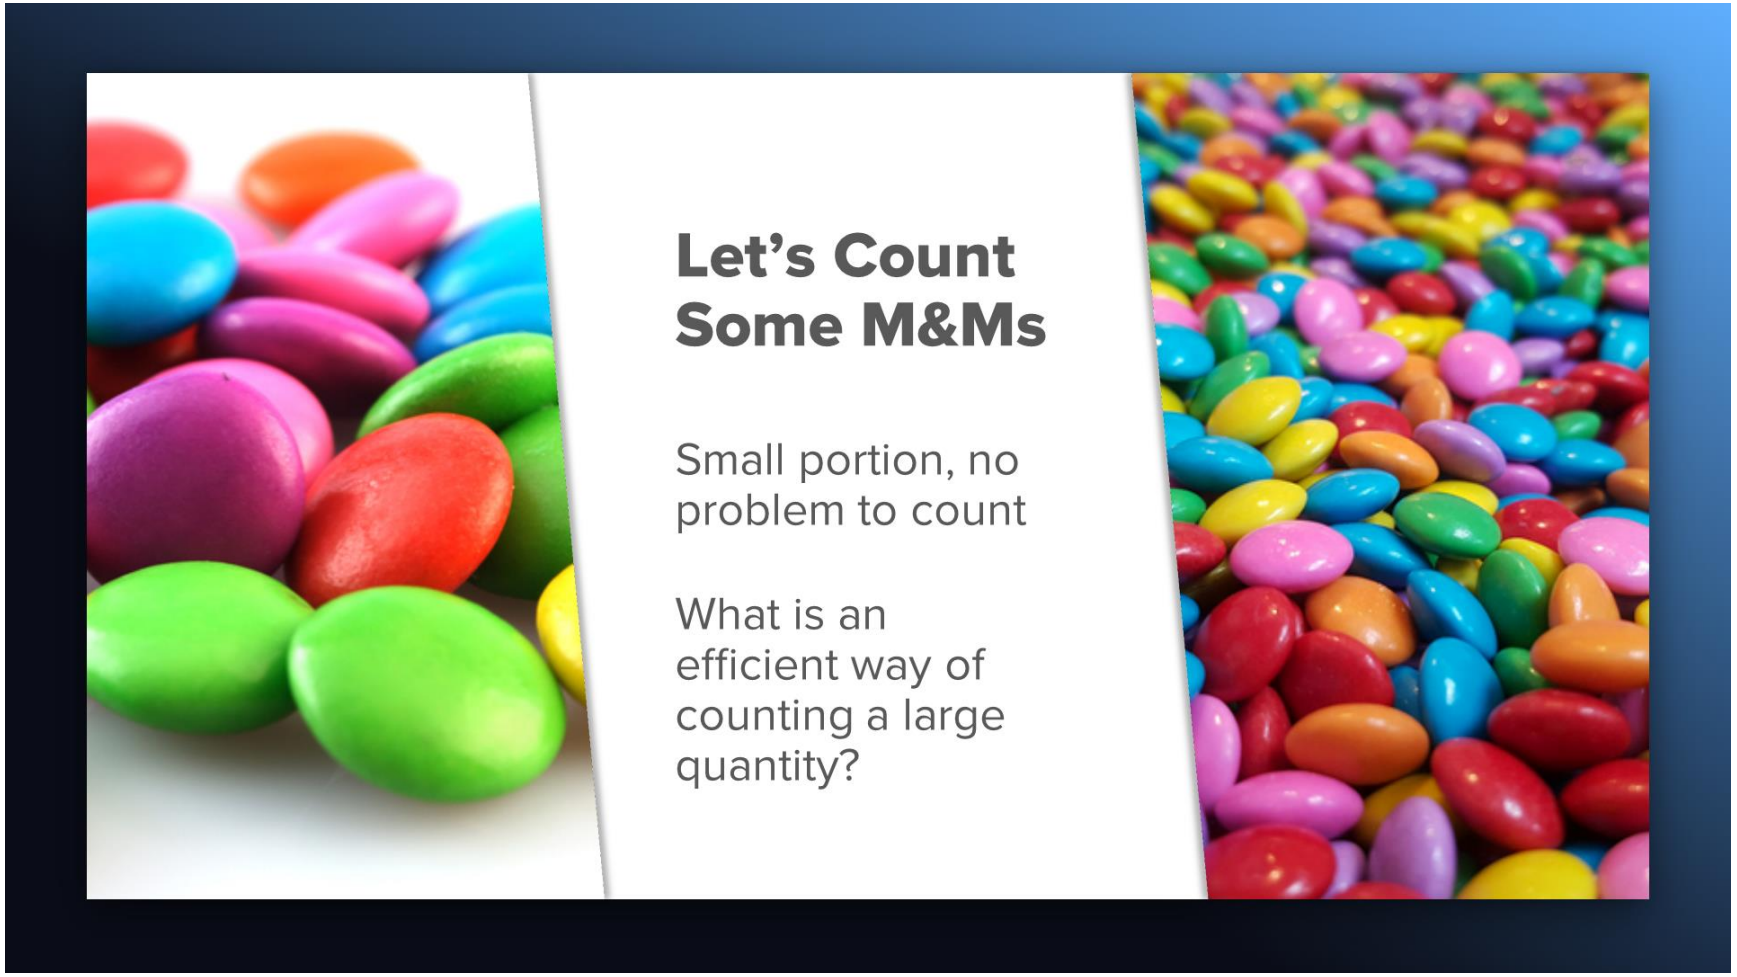

Copyright ©The Regents of the University of California

Slide 12: Measure by Weight?

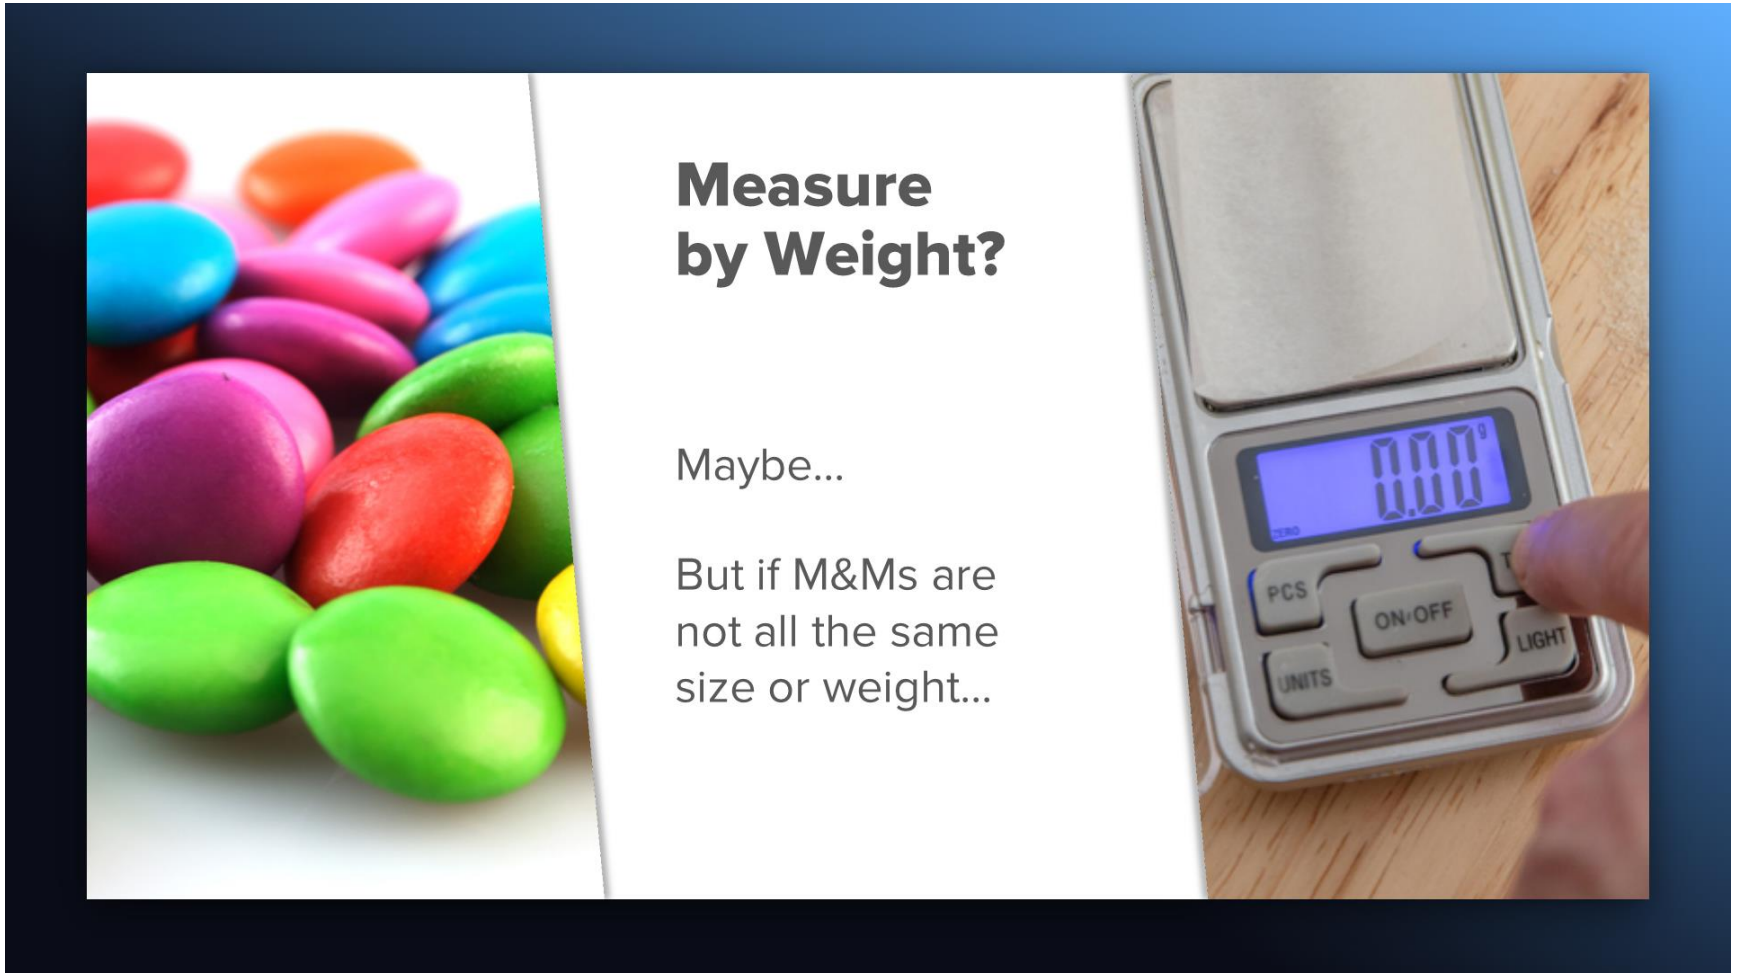

Copyright ©The Regents of the University of California

Slide 13: Involving Others

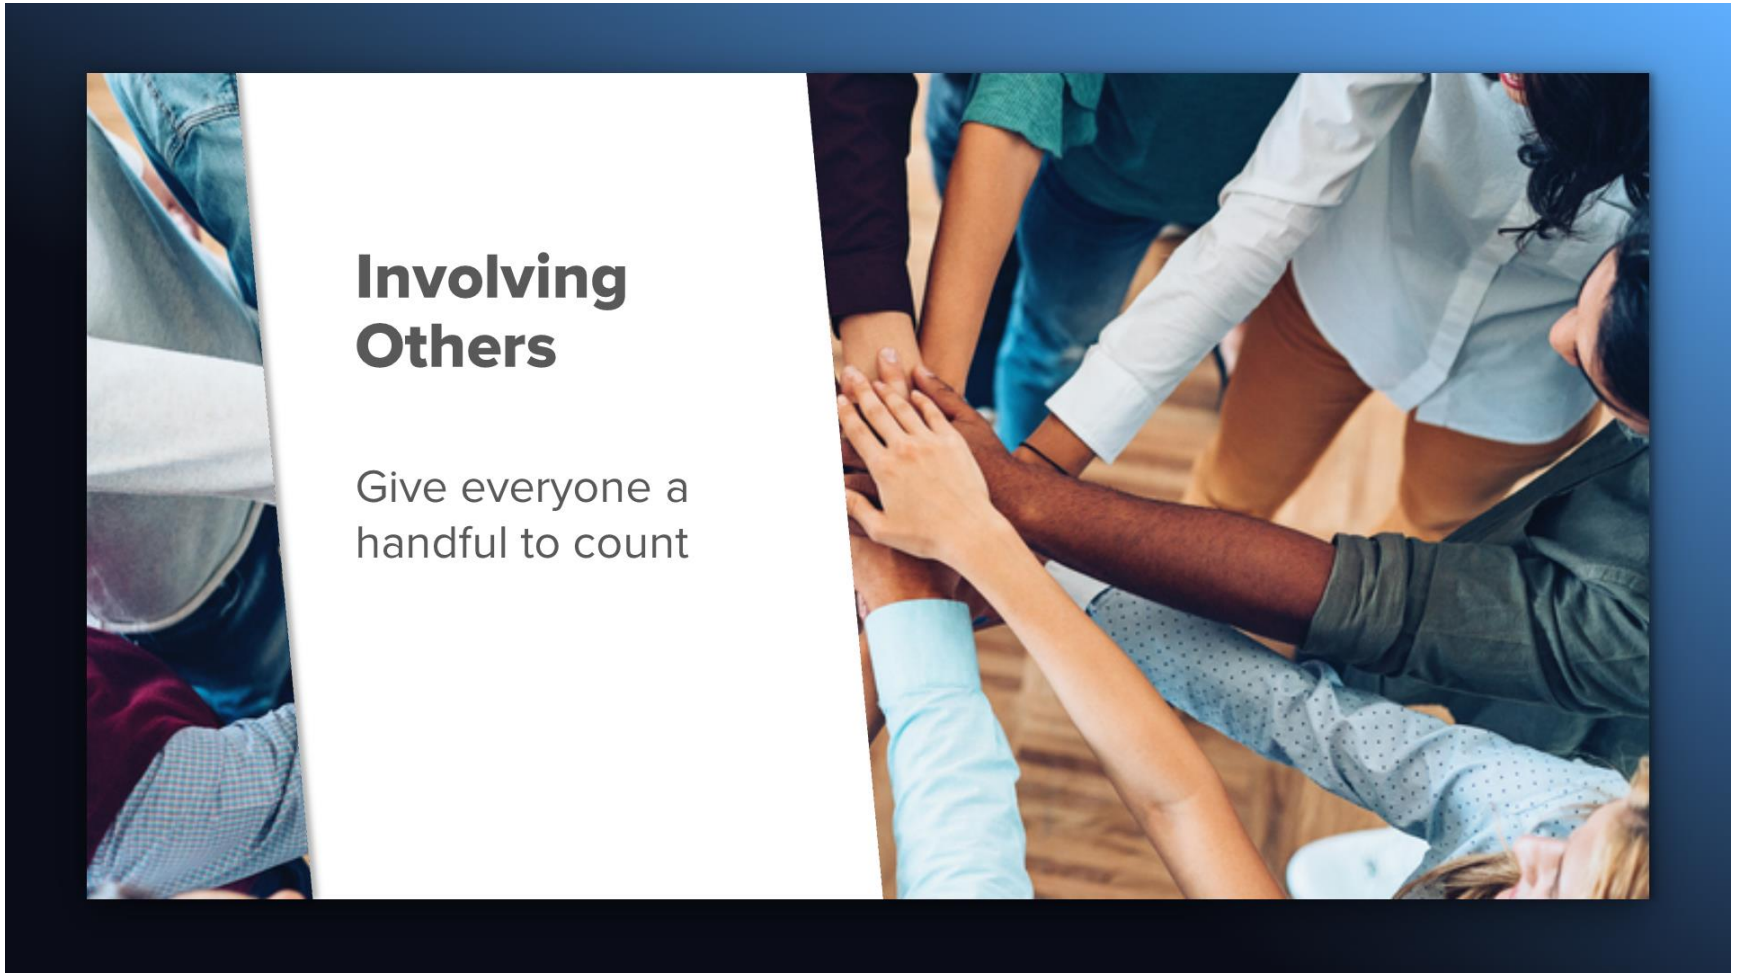

Copyright ©The Regents of the University of California

Slide 14: How Do You Aggregate All of the Counts Among Your Friends?

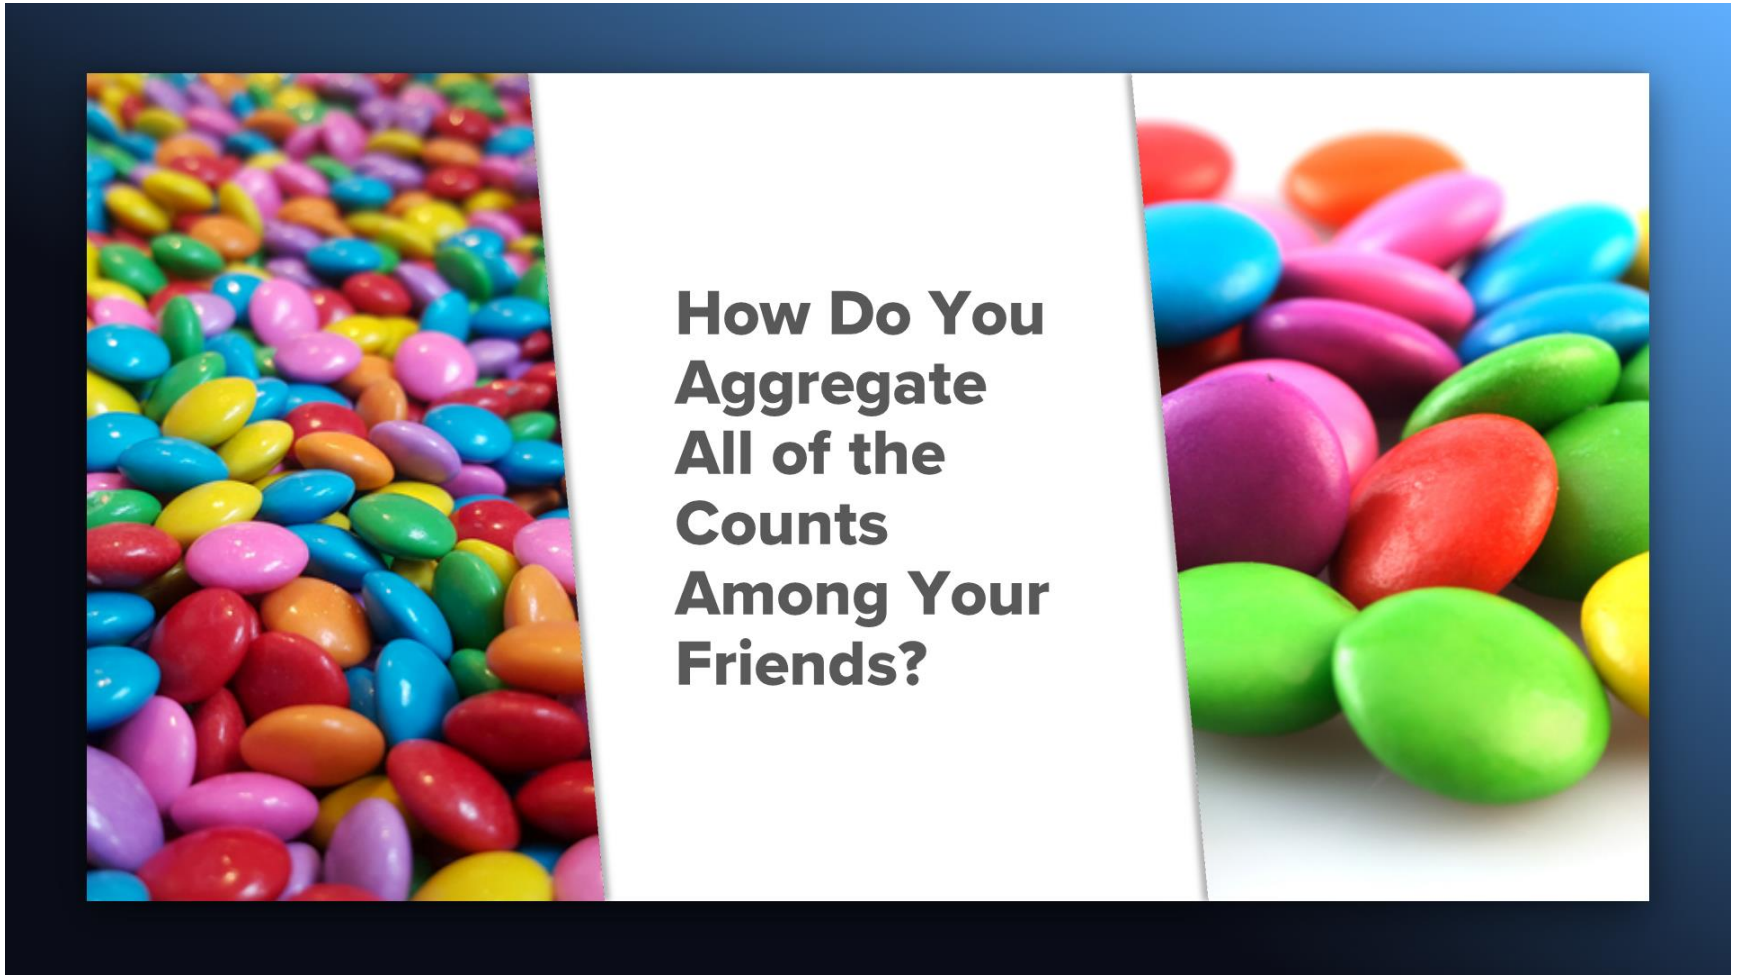

Copyright ©The Regents of the University of California

Slide 15: Drivers and Executors

Copyright ©The Regents of the University of California

Slide 16: Drivers and Executors

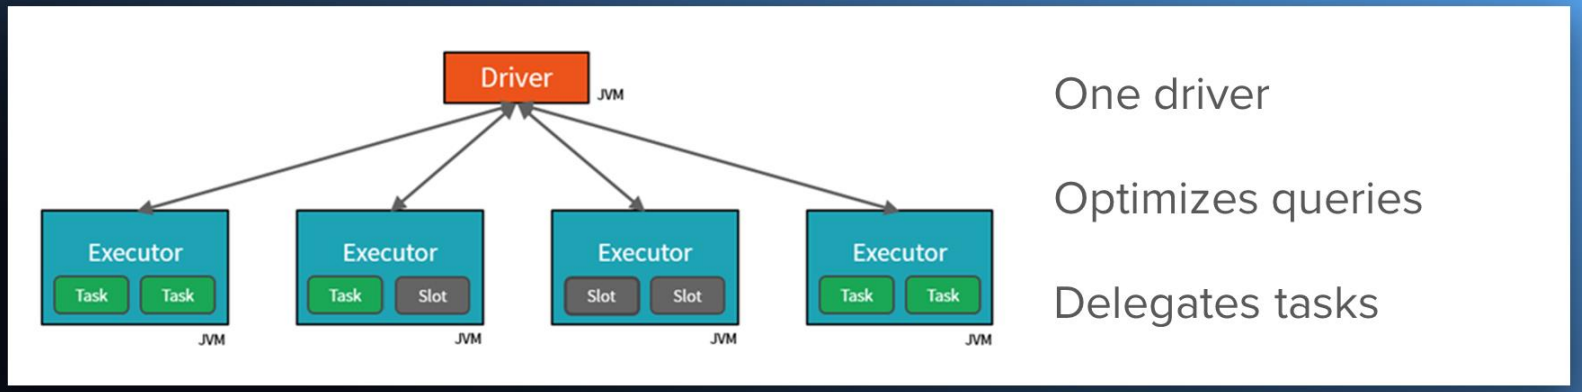

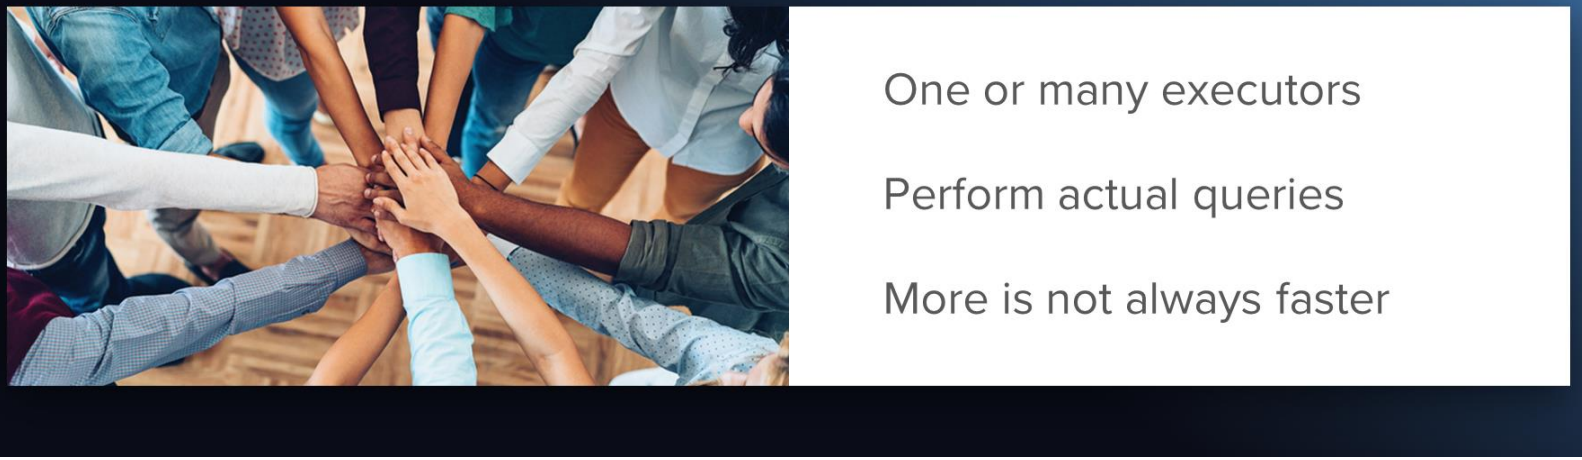

Copyright ©The Regents of the University of California

Slide 17: Why More Computing Power Isn’t Always Faster

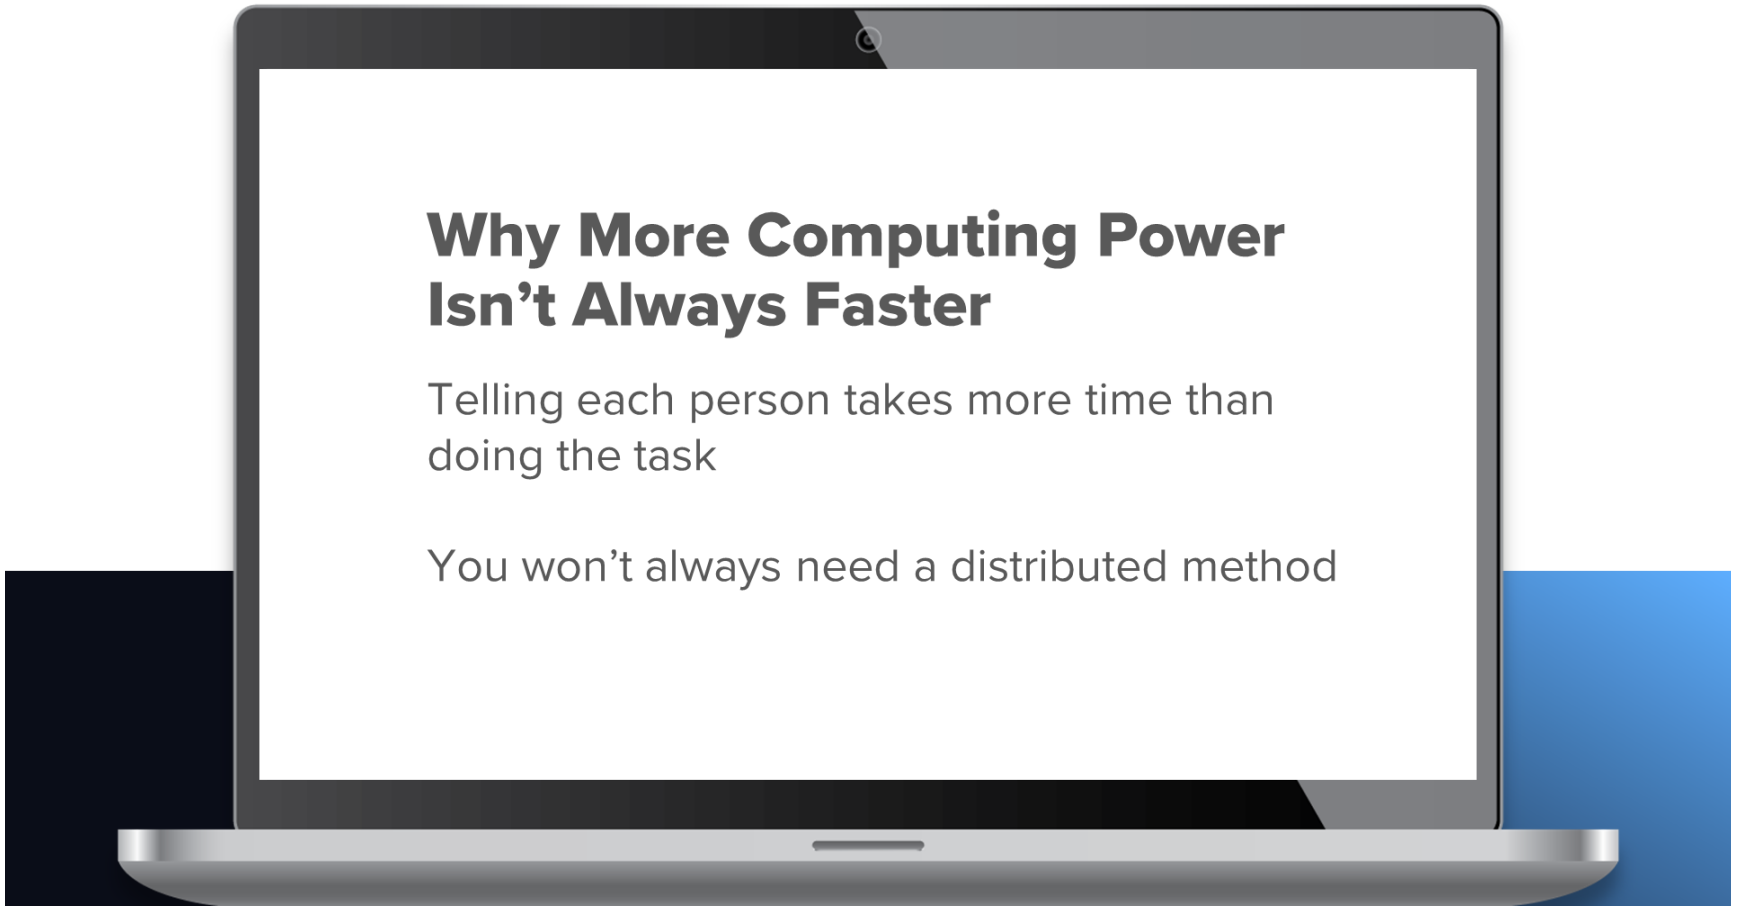

Copyright ©The Regents of the University of California

Slide 18: Distributing Computation is Parallelism

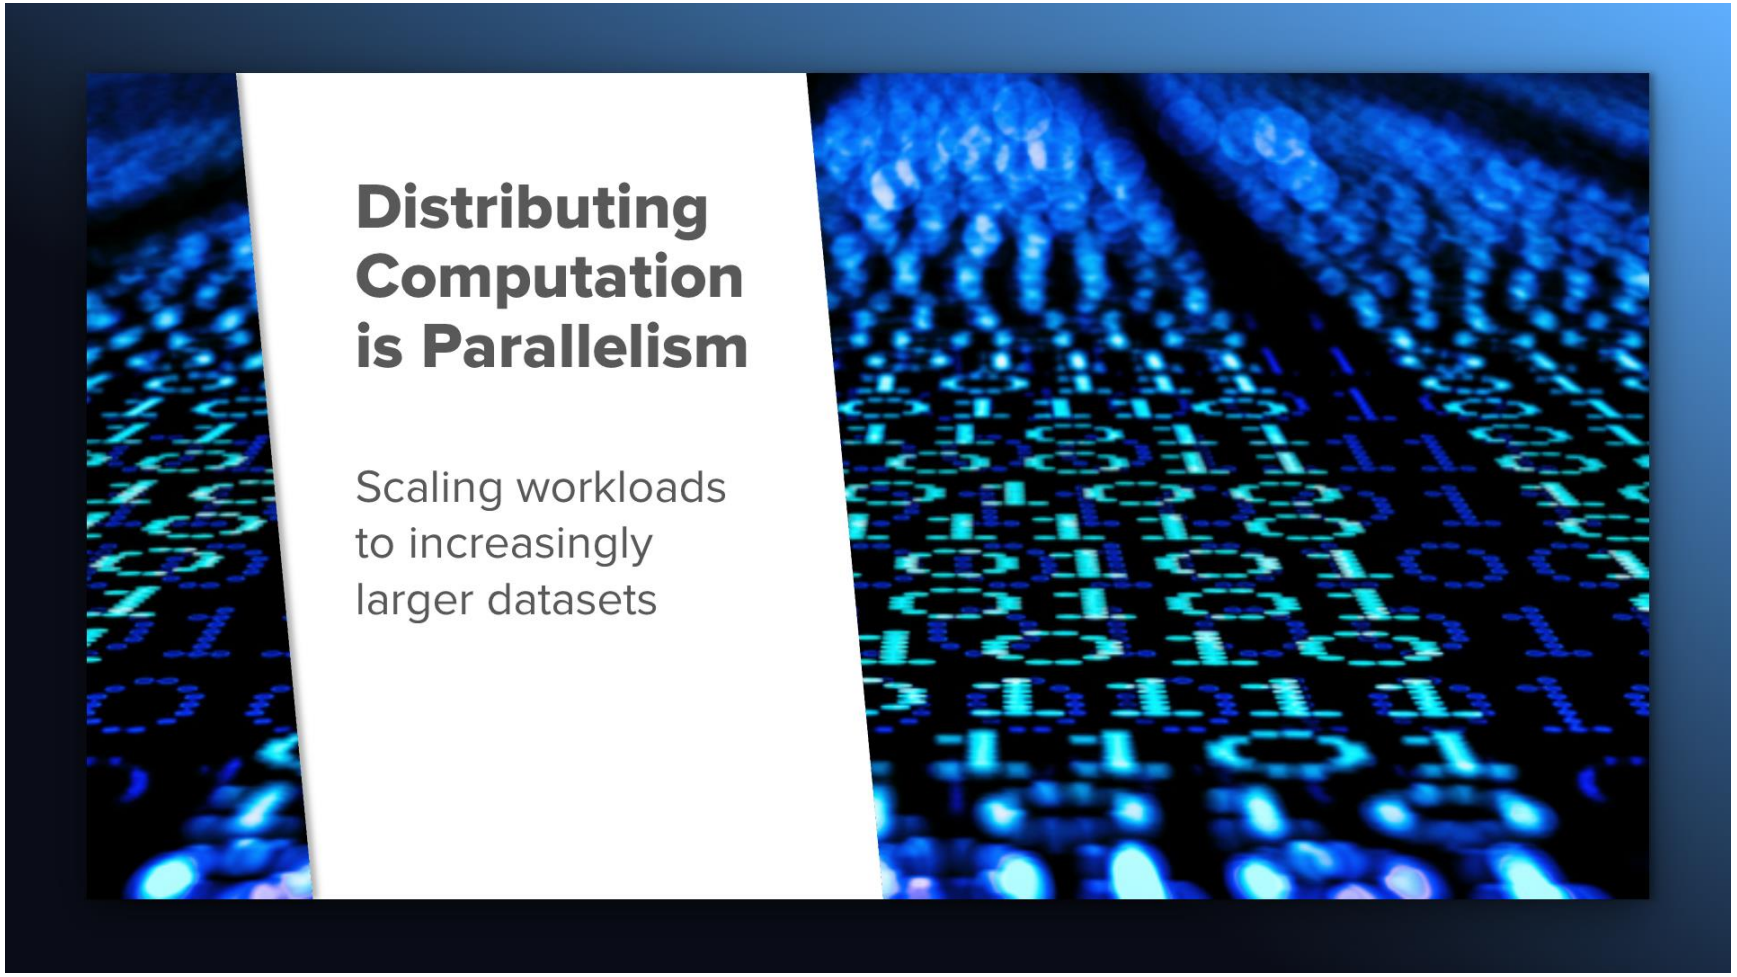

Copyright ©The Regents of the University of California

Slide 19: Amdahl's Law

Amdahl's Law

The amount of acceleration we would see from parallelizing a task is a function of what portion of the task can be completed in parallel

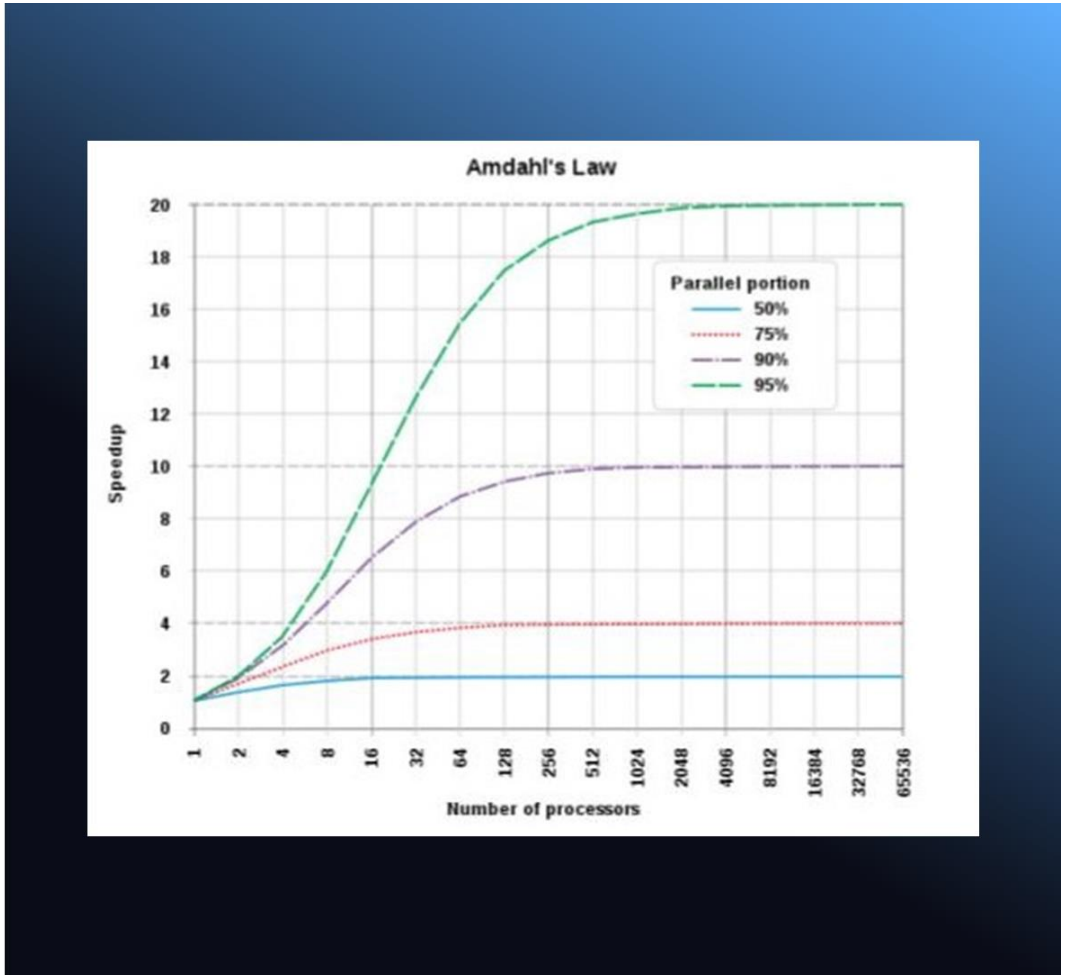

Copyright ©The Regents of the University of California

Slide 20: Linear Scalability

Copyright ©The Regents of the University of California

Slide 21: Scalability

Scale out

More data to process than on one machine

Speed up

More computer resources may speed up your query

Copyright ©The Regents of the University of California

Slide 22: Coming Up

Coming Up

Learn about core Spark concepts

Copyright ©The Regents of the University of California